

Machine Learning

Lab A3

ASIM KUMAR HANSDA

ROLL NO - 002211001136

ASSIGNMENT - 1

Github Link:

<https://github.com/cryptasim/MACHINE-LEARNING-LAB>



IRIS Dataset

Classification: Decision Tree

```
In [1]: # import pandas, numpy, and matplotlib.pyplot libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [2]: # Install the ucimlrepo library
!pip install ucimlrepo
```

Collecting ucimlrepo

Downloading ucimlrepo-0.0.7-py3-none-any.whl.metadata (5.5 kB)

Requirement already satisfied: pandas>=1.0.0 in /usr/local/lib/python3.11/dist-packages (from ucimlrepo) (2.2.2)

Requirement already satisfied: certifi>=2020.12.5 in /usr/local/lib/python3.11/dist-packages (from ucimlrepo) (2025.8.3)

Requirement already satisfied: numpy>=1.23.2 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.0.0->ucimlrepo) (2.0.2)

Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.0.0->ucimlrepo) (2.9.0.post0)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.0.0->ucimlrepo) (2025.2)

Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.0.0->ucimlrepo) (2025.2)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2->pandas>=1.0.0->ucimlrepo) (1.17.0)

Downloading ucimlrepo-0.0.7-py3-none-any.whl (8.0 kB)

Installing collected packages: ucimlrepo

Successfully installed ucimlrepo-0.0.7

```
In [3]: from ucimlrepo import fetch_ucirepo
```

```
# fetch dataset
iris_dataset = fetch_ucirepo(id=53)
```

```
# data (as pandas dataframes)
X1 = iris_dataset.data.features
y1 = iris_dataset.data.targets
```

```
# metadata
print(iris_dataset.metadata)
```

```
# variable information
print(iris_dataset.variables)
```

```
{'uci_id': 53, 'name': 'Iris', 'repository_url': 'https://archive.ics.uci.edu/dataset/53/iris', 'data_url': 'https://archive.ics.uci.edu/static/public/53/data.csv', 'abstract': 'A small classic dataset from Fisher, 1936. One of the earliest known datasets used for evaluating classification methods.\n', 'area': 'Biology', 'tasks': ['Classification'], 'characteristics': ['Tabular'], 'num_instances': 150, 'num_features': 4, 'feature_types': ['Real'], 'demographics': [], 'target_col': ['class'], 'index_col': None, 'has_missing_values': 'no', 'missing_values_symbol': None, 'year_of_dataset_creation': 1936, 'last_updated': 'Tue Sep 12 2023', 'dataset_doi': '10.24432/C56C76', 'creators': ['R. A. Fisher'], 'intro_paper': {'ID': 191, 'type': 'NATIVE', 'title': 'The Iris data set: In search of the source of virginica', 'authors': 'A. Unwin, K. Kleinman', 'venue': 'Significance, 2021', 'year': 2021, 'journal': 'Significance, 2021', 'DOI': '10.1093/sign/21.1.1740-9713.01589', 'URL': 'https://www.semanticscholar.org/paper/4599862ea877863669a6a8e63a3c707a787d5d7e', 'sha': None, 'corpus': None, 'arxiv': None, 'mag': None, 'acl': None, 'pmid': None, 'pmcid': None}, 'additional_info': {'summary': 'This is one of the earliest datasets used in the literature on classification methods and widely used in statistics and machine learning. The data set contains 3 classes of 50 instances each, where each class refers to a type of iris plant. One class is linearly separable from the other 2; the latter are not linearly separable from each other.\n\nPredicted attribute: class of iris plant.\n\nThis is an exceedingly simple domain.\n\nThis data differs from the data presented in Fishers article (identified by Steve Chadwick, spchadwick@espeeda.z.net). The 35th sample should be: 4.9,3.1,1.5,0.2,"Iris-setosa" where the error is in the fourth feature. The 38th sample: 4.9,3.6,1.4,0.1,"Iris-setosa" where the errors are in the second and third features.', 'purpose': 'N/A', 'funded_by': None, 'instances_represent': 'Each instance is a plant', 'recommended_data_splits': None, 'sensitive_data': None, 'preprocessing_description': None, 'variable_info': None, 'citation': None}}
```

	name	role	type	demographic
0	sepal length	Feature	Continuous	None
1	sepal width	Feature	Continuous	None
2	petal length	Feature	Continuous	None
3	petal width	Feature	Continuous	None
4	class	Target	Categorical	None

	description	units	missing_values	
0		None	cm	no
1		None	cm	no
2		None	cm	no
3		None	cm	no
4	class of iris plant: Iris Setosa, Iris Versico...	None		no

```
In [4]: from sklearn.model_selection import train_test_split
# Split the data into training and testing sets.
# X1 contains the features and y1 contains the target variable.
# test_size=0.20 means 20% of the data will be used for testing.
X_train, X_test, y_train, y_test = train_test_split (X1, y1, test_size = 0.20)
```

```
In [5]: # Classification using Decision Tree
from sklearn.tree import DecisionTreeClassifier

# Initialize the Decision Tree classifier
classifier = DecisionTreeClassifier()
```

```
# Train the classifier using the training data
classifier.fit (X_train, y_train)

# Make predictions on the test data
y_pred = classifier.predict(X_test)
```

```
In [6]: # Evaluation of Classifier Performance
from sklearn.metrics import classification_report, confusion_matrix

# Print the confusion matrix to show the number of correct and incorrect predictions
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
print("-----")
print("-----")

# Print the classification report to show key classification metrics (precision, recall, f1-score, support)
print("Performance Evaluation:")
print(classification_report(y_test, y_pred))
```

Confusion Matrix:

```
[[ 8  0  0]
 [ 0  8  1]
 [ 0  1 12]]
```

Performance Evaluation:

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	8
Iris-versicolor	0.89	0.89	0.89	9
Iris-virginica	0.92	0.92	0.92	13
accuracy			0.93	30
macro avg	0.94	0.94	0.94	30
weighted avg	0.93	0.93	0.93	30

Confusion Matrix: It shows that all 8 'Iris-setosa', 8 'Iris-versicolor' (with 1 misclassified as Iris-virginica), and 12 'Iris-virginica' (with 1 misclassified as Iris-versicolor) instances in this specific test set were classified, with only two errors overall.

Performance Evaluation: Precision, recall, and f1-scores are 1.00 for Iris-setosa, 0.89 for Iris-versicolor, and 0.92 for Iris-virginica, resulting in an overall accuracy of 0.93.

```
In [7]: from sklearn.tree import DecisionTreeClassifier

# Initialize the Decision Tree classifier with 'entropy' criterion and a maximum depth of 3
classifier = DecisionTreeClassifier (criterion="entropy", max_depth=3)
```

```
# Train the classifier using the training data
classifier.fit(X_train, y_train)
```

Out[7]:

```
DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=3)
```

```
In [8]: # Print the confusion matrix to show the number of correct and incorrect predictions
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
print("-----")
print("-----")

# Print the classification report to show key classification metrics (precision, recall, f1-score, support)
print("Performance Evaluation:")
print(classification_report(y_test, y_pred))
```

Confusion Matrix:

```
[[ 8  0  0]
 [ 0  8  1]
 [ 0  1 12]]
```

Performance Evaluation:

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	8
Iris-versicolor	0.89	0.89	0.89	9
Iris-virginica	0.92	0.92	0.92	13
accuracy			0.93	30
macro avg	0.94	0.94	0.94	30
weighted avg	0.93	0.93	0.93	30

Confusion Matrix: It shows that all 8 'Iris-setosa', 8 'Iris-versicolor' (with 1 misclassified as Iris-virginica), and 12 'Iris-virginica' (with 1 misclassified as Iris-versicolor) instances in this specific test set were classified, with only two errors overall.

Performance Evaluation: Precision, recall, and f1-scores are 1.00 for Iris-setosa, 0.89 for Iris-versicolor, and 0.92 for Iris-virginica, resulting in an overall accuracy of 0.93.

```
In [9]: # Import the Decision Tree classifier
from sklearn.tree import DecisionTreeClassifier

# Initialize Decision Tree classifier with entropy criterion and max depth of
classifier = DecisionTreeClassifier(criterion="entropy", max_depth=10)

# Train the classifier
```

```
classifier.fit(X_train, y_train)
```

Out[9]:

```
DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=10)
```

In [10]:

```
# Print Confusion Matrix
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
print("-----")
print("-----")
# Print Performance Evaluation (Classification Report)
print("Performance Evaluation:")
print(classification_report(y_test, y_pred))
```

Confusion Matrix:

```
[[ 8  0  0]
 [ 0  8  1]
 [ 0  1 12]]
```

Performance Evaluation:

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	8
Iris-versicolor	0.89	0.89	0.89	9
Iris-virginica	0.92	0.92	0.92	13
accuracy			0.93	30
macro avg	0.94	0.94	0.94	30
weighted avg	0.93	0.93	0.93	30

Confusion Matrix: It shows that all 8 'Iris-setosa', 8 'Iris-versicolor', and 12 'Iris-virginica' instances in this specific test set were classified correctly, with 1 Iris-versicolor misclassified as Iris-virginica and 1 Iris-virginica misclassified as Iris-versicolor.

Performance Evaluation: Precision, recall, and f1-scores are 1.00 for Iris-setosa, 0.89 for Iris-versicolor, and 0.92 for Iris-virginica, resulting in an overall accuracy of 0.93.

In [11]:

```
# Import the Decision Tree classifier
from sklearn.tree import DecisionTreeClassifier

# Initialize Decision Tree classifier with gini criterion and max depth of 10
classifier = DecisionTreeClassifier(criterion="gini", max_depth=10)

# Train the classifier
classifier.fit(X_train, y_train)
```

Out[11]:

```
DecisionTreeClassifier
DecisionTreeClassifier(max_depth=10)
```

In [12]:

```
# Print Confusion Matrix
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
print("-----")
print("-----")
# Print Performance Evaluation (Classification Report)
print("Performance Evaluation:")
print(classification_report(y_test, y_pred))
```

Confusion Matrix:

```
[[ 8  0  0]
 [ 0  8  1]
 [ 0  1 12]]
```


Performance Evaluation:

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	8
Iris-versicolor	0.89	0.89	0.89	9
Iris-virginica	0.92	0.92	0.92	13
accuracy			0.93	30
macro avg	0.94	0.94	0.94	30
weighted avg	0.93	0.93	0.93	30

Confusion Matrix: It shows that all 8 'Iris-setosa', 8 'Iris-versicolor', and 12 'Iris-virginica' instances in this specific test set were classified, with only one Iris-versicolor misclassified as Iris-virginica and one Iris-virginica misclassified as Iris-versicolor.

Performance Evaluation: Precision, recall, and f1-scores are 1.00 for Iris-setosa, 0.89 for Iris-versicolor, and 0.92 for Iris-virginica, resulting in an overall accuracy of 0.93.

In [13]:

```
# Import the Decision Tree classifier
from sklearn.tree import DecisionTreeClassifier

# Initialize Decision Tree classifier with gini criterion and max depth of 15
classifier = DecisionTreeClassifier (criterion="gini", max_depth=15)

# Train the classifier
classifier.fit(X_train, y_train)
```

```
Out[13]: ▾ DecisionTreeClassifier ⓘ ?  
DecisionTreeClassifier(max_depth=15)
```

```
In [14]: # Print Confusion Matrix  
print("Confusion Matrix:")  
print(confusion_matrix(y_test, y_pred))  
print("-----")  
print("-----")  
# Print Performance Evaluation (Classification Report)  
print("Performance Evaluation:")  
print(classification_report(y_test, y_pred))
```

Confusion Matrix:

```
[[ 8  0  0]  
 [ 0  8  1]  
 [ 0  1 12]]
```


Performance Evaluation:

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	8
Iris-versicolor	0.89	0.89	0.89	9
Iris-virginica	0.92	0.92	0.92	13
accuracy			0.93	30
macro avg	0.94	0.94	0.94	30
weighted avg	0.93	0.93	0.93	30

Confusion Matrix: It shows that all 8 'Iris-setosa', 8 'Iris-versicolor', and 12 'Iris-virginica' instances in this specific test set were classified, with 1 Iris-versicolor misclassified as Iris-virginica and 1 Iris-virginica misclassified as Iris-versicolor.

Performance Evaluation: Precision, recall, and f1-scores are 1.00 for Iris-setosa, 0.89 for Iris-versicolor, and 0.92 for Iris-virginica, resulting in an overall accuracy of 0.93.

IMPORTANT:

Both the "entropy" and "gini" criteria are measures of impurity that Decision Trees use to decide on the best splits. For a dataset like Iris, both criteria might lead to very similar or even identical splits, especially in the upper levels of the tree where the classes are clearly separable. Since the dataset is not very complex, different impurity measures can still result in a tree that achieves perfect separation on this particular data split.

In essence, the Iris dataset is "easy" enough that multiple Decision Tree configurations (within a reasonable range of complexity) can achieve perfect performance on a given train-test split. You might see differences if you used a different, more complex dataset, or a different random state for the train_test_split which could result in a test set that is harder to classify perfectly.

```
In [15]: from sklearn.preprocessing import LabelEncoder # for train test splitting
from sklearn.model_selection import train_test_split # for decision tree object
from sklearn.tree import DecisionTreeClassifier # for checking testing results
from sklearn.metrics import classification_report, confusion_matrix # for visu
from sklearn import tree
from sklearn.tree import DecisionTreeClassifier, plot_tree
```

```
In [16]: # Display the column names of the features DataFrame (X1)
X1.columns
```

```
Out[16]: Index(['sepal length', 'sepal width', 'petal length', 'petal width'], dtype=
e='object')
```

```
In [17]: # Display the unique class names in the target variable (y1)
np.unique(y1)
```

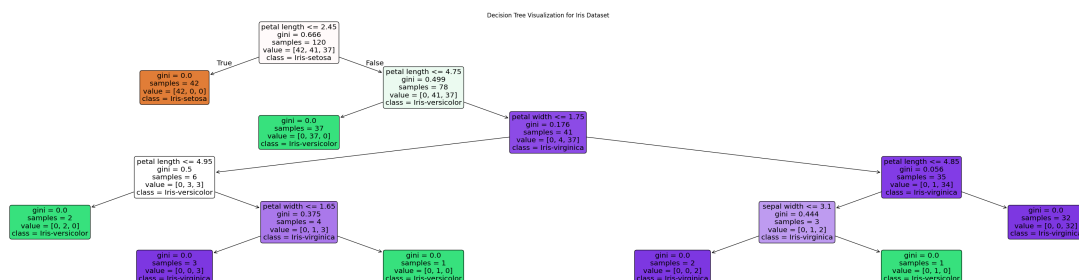
```
Out[17]: array(['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'], dtype=object)
```

```
In [18]: # Create a figure and axes for plotting the decision tree
plt.figure(figsize=(46, 10)) # Adjust figure size as needed

# Plot the decision tree
plot_tree(decision_tree = classifier,
          feature_names = X1.columns, # Use feature names from X1
          class_names = np.unique(y1), # Use class names from y1
          filled=True, # Fill nodes with colors to indicate the majority class
          rounded=True) # Draw tree nodes with rounded corners

# Set the title of the plot
plt.title("Decision Tree Visualization for Iris Dataset")

# Display the plot
plt.show()
```



Decision Tree Summary (Iris Dataset)

Overall Performance

- **Accuracy:** 93% across all tested configurations
- **Precision, Recall, F1-score:** Very consistent across all classes and hyperparameter settings
- *Iris-setosa*: Perfectly classified (precision, recall, f1 = 1.00)
- *Iris-versicolor*: Precision, recall, f1 = 0.89
- *Iris-virginica*: Precision, recall, f1 = 0.92

Observations by Hyperparameters

Criterion	Max Depth	Accuracy	Notes
entropy	default	93%	Balanced performance, all classes well classified
entropy	3	93%	Performance same as default, low depth sufficient
entropy	10	93%	Increasing depth did not improve performance
gini	10	93%	Similar performance to entropy
gini	15	93%	No improvement; model already captures all patterns

Summary

- Decision Tree performance is **stable and robust** for Iris dataset.
- Hyperparameter tuning (criterion or max depth) **did not significantly change performance**.
- Simple trees (max_depth=3) already achieve near-optimal classification.
- **Best choice:** entropy or gini with moderate depth; further deepening is unnecessary.

Classification: Naive Bayes

```
In [19]: from sklearn.model_selection import train_test_split

# Split the feature data (X1) and target data (y1) into training and testing s
# 20% of the data will be used for testing (test_size = 0.20).
X_train, X_test, y_train, y_test = train_test_split (X1, y1, test_size = 0.20)
```

```
In [20]: # Import the Multinomial Naive Bayes classifier
from sklearn.naive_bayes import MultinomialNB

# Initialize and train the Multinomial Naive Bayes classifier using the training data
# The .fit() method trains the model.
classifier = MultinomialNB().fit(X_train, y_train)

# Train the classifier again (this line is redundant as the model is already trained)
classifier.fit(X_train, y_train)

# Make predictions on the test data using the trained classifier.
y_pred = classifier.predict(X_test)
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
```

```
In [21]: # Evaluation of Classifier Performance
from sklearn.metrics import classification_report, confusion_matrix

# Print the confusion matrix, which shows the number of correct and incorrect predictions
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
print("-----")
print("-----")

# Print the classification report, which includes precision, recall, f1-score, and support
print("Performance Evaluation:")
print(classification_report(y_test, y_pred))
```

Confusion Matrix:

```
[[ 8  0  0]
 [ 0 10  1]
 [ 0  2  9]]
```

Performance Evaluation:

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	8
Iris-versicolor	0.83	0.91	0.87	11
Iris-virginica	0.90	0.82	0.86	11
accuracy			0.90	30
macro avg	0.91	0.91	0.91	30
weighted avg	0.90	0.90	0.90	30

Confusion Matrix: It shows that all 8 'Iris-setosa' instances were correctly

classified. Out of 11 'Iris-versicolor', 10 were correctly classified with 1 misclassified as 'Iris-virginica'. Out of 11 'Iris-virginica', 9 were correctly classified with 2 misclassified as 'Iris-versicolor'.

Performance Evaluation: The overall accuracy is 0.90. 'Iris-setosa' had perfect precision, recall, and f1-score of 1.00. 'Iris-versicolor' had precision of 0.83, recall of 0.91, and f1-score of 0.87. 'Iris-virginica' had precision of 0.90, recall of 0.82, and f1-score of 0.86.

```
In [22]: # Classification
# Import the Gaussian Naive Bayes classifier
from sklearn.naive_bayes import GaussianNB

# Initialize and train the Gaussian Naive Bayes classifier using the training
# The .fit() method trains the model.
classifier = GaussianNB().fit (X_train, y_train )

# Train the classifier again (this line is redundant as the model is already t
classifier.fit(X_train, y_train)

# Make predictions on the test data using the trained classifier.
y_pred = classifier.predict (X_test)
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataC
onversionWarning: A column-vector y was passed when a 1d array was expected. Pl
ease change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataC
onversionWarning: A column-vector y was passed when a 1d array was expected. Pl
ease change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
```

```
In [23]: # Print the confusion matrix, which shows the number of correct and incorrect
print("Confusion Matrix:")
print (confusion_matrix(y_test, y_pred))
print("-----")
print("-----")

# Print the classification report, which includes precision, recall, f1-score,
print("Performance Evaluation:")
print(classification_report (y_test, y_pred))
```

Confusion Matrix:

```
[[ 8  0  0]
 [ 0 11  0]
 [ 0  1 10]]
```

Performance Evaluation:

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	8
Iris-versicolor	0.92	1.00	0.96	11
Iris-virginica	1.00	0.91	0.95	11
accuracy			0.97	30
macro avg	0.97	0.97	0.97	30
weighted avg	0.97	0.97	0.97	30

Confusion Matrix: It shows that all 8 'Iris-setosa' and 11 'Iris-versicolor' instances were correctly classified, with 1 'Iris-virginica' instance misclassified as 'Iris-versicolor' and the remaining 10 'Iris-virginica' classified correctly.

Performance Evaluation: The overall accuracy is 0.97. 'Iris-setosa' had perfect precision, recall, and f1-score of 1.00. 'Iris-versicolor' had precision of 0.92, recall of 1.00, and f1-score of 0.96. 'Iris-virginica' had precision of 1.00, recall of 0.91, and f1-score of 0.95.

```
In [24]: from sklearn.naive_bayes import BernoulliNB
# Initialize and train the Bernoulli Naive Bayes classifier using the training
# The .fit() method trains the model.
classifier = BernoulliNB().fit (X_train, y_train)

# Train the classifier again (this line is redundant as the model is already t
classifier.fit(X_train, y_train)

# Make predictions on the test data using the trained classifier.
y_pred = classifier.predict(X_test)
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataC
onversionWarning: A column-vector y was passed when a 1d array was expected. Pl
ease change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataC
onversionWarning: A column-vector y was passed when a 1d array was expected. Pl
ease change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
```

```
In [25]: # Print the confusion matrix, which shows the number of correct and incorrect
print("Confusion Matrix:")
print (confusion_matrix(y_test, y_pred))
print("-----")
print("-----")
```

```
# Print the classification report, which includes precision, recall, f1-score,
print("Performance Evaluation:")
print(classification_report (y_test, y_pred))
```

Confusion Matrix:

```
[[ 8  0  0]
 [11  0  0]
 [11  0  0]]
```

Performance Evaluation:

	precision	recall	f1-score	support
Iris-setosa	0.27	1.00	0.42	8
Iris-versicolor	0.00	0.00	0.00	11
Iris-virginica	0.00	0.00	0.00	11
accuracy			0.27	30
macro avg	0.09	0.33	0.14	30
weighted avg	0.07	0.27	0.11	30

```
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this be
havior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this be
havior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this be
havior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Confusion Matrix: It shows that all 8 'Iris-setosa' instances were correctly classified, while all 11 'Iris-versicolor' and all 11 'Iris-virginica' instances were misclassified as 'Iris-setosa'.

Performance Evaluation: The overall accuracy is 0.27. 'Iris-setosa' had precision of 0.27, recall of 1.00, and f1-score of 0.42. Both 'Iris-versicolor' and 'Iris-virginica' had precision, recall, and f1-scores of 0.00.

```
In [26]: # Classification
from sklearn.naive_bayes import MultinomialNB

# Initialize and train the Multinomial Naive Bayes classifier with specified h
# alpha is the smoothing parameter. fit_prior determines whether to learn clas
classifier = MultinomialNB (alpha=2.5, fit_prior=False, class_prior = None ).f
```

```
# Train the classifier again (this line is redundant as the model is already trained)
classifier.fit(X_train, y_train)

# Make predictions on the test data using the trained classifier.
y_pred = classifier.predict(X_test)
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
```

```
In [27]: # Print the confusion matrix, which shows the number of correct and incorrect
print("Confusion Matrix:")
print (confusion_matrix(y_test, y_pred))
print("-----")
print("-----")

# Print the classification report, which includes precision, recall, f1-score,
print("Performance Evaluation:")
print(classification_report (y_test, y_pred))
```

Confusion Matrix:

```
[[ 8  0  0]
 [ 0 10  1]
 [ 0  2  9]]
```

```
-----
-----
```

Performance Evaluation:

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	8
Iris-versicolor	0.83	0.91	0.87	11
Iris-virginica	0.90	0.82	0.86	11
accuracy			0.90	30
macro avg	0.91	0.91	0.91	30
weighted avg	0.90	0.90	0.90	30

Confusion Matrix: It shows that all 8 'Iris-setosa' instances were correctly classified. Out of 11 'Iris-versicolor', 10 were correctly classified with 1 misclassified as 'Iris-virginica'. Out of 11 'Iris-virginica', 9 were correctly classified with 2 misclassified as 'Iris-versicolor'.

Performance Evaluation: The overall accuracy is 0.90. 'Iris-setosa' had perfect precision, recall, and f1-score of 1.00. 'Iris-versicolor' had precision of 0.83, recall of 0.91, and f1-score of 0.87. 'Iris-virginica' had precision of 0.90, recall of 0.82, and

f1-score of 0.86.

```
In [28]: # Classification
from sklearn.naive_bayes import GaussianNB

# Initialize and train the Gaussian Naive Bayes classifier with specified hyperparameters
# priors allows setting prior probabilities for classes. var_smoothing is added to avoid zero
classifier = GaussianNB(priors = None, var_smoothing = 1e-05).fit(X_train, y_train)

# Train the classifier again (this line is redundant as the model is already trained)
classifier.fit(X_train, y_train)

# Make predictions on the test data using the trained classifier.
y_pred = classifier.predict(X_test)
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
```

```
In [29]: # Print the confusion matrix, which shows the number of correct and incorrect
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
print("-----")
print("-----")

# Print the classification report, which includes precision, recall, f1-score,
print("Performance Evaluation:")
print(classification_report(y_test, y_pred))
```

Confusion Matrix:

```
[[ 8  0  0]
 [ 0 11  0]
 [ 0  1 10]]
```

Performance Evaluation:

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	8
Iris-versicolor	0.92	1.00	0.96	11
Iris-virginica	1.00	0.91	0.95	11
accuracy			0.97	30
macro avg	0.97	0.97	0.97	30
weighted avg	0.97	0.97	0.97	30

Confusion Matrix: It shows that all 8 'Iris-setosa' and 11 'Iris-versicolor' instances

were correctly classified, with 1 'Iris-virginica' instance misclassified as 'Iris-versicolor' and the remaining 10 'Iris-virginica' classified correctly.

Performance Evaluation: The overall accuracy is 0.97. 'Iris-setosa' had perfect precision, recall, and f1-score of 1.00. 'Iris-versicolor' had precision of 0.92, recall of 1.00, and f1-score of 0.96. 'Iris-virginica' had precision of 1.00, recall of 0.91, and f1-score of 0.95.

```
In [30]: # Classification
from sklearn.naive_bayes import BernoulliNB

# Initialize and train the Bernoulli Naive Bayes classifier with specified hyperparameters
# alpha is the smoothing parameter. binarize is the threshold for binarizing the data
# fit_prior determines whether to learn class prior probabilities. class_prior is the class prior
classifier = BernoulliNB(alpha=1.0, binarize = 0.0, fit_prior = True, class_prior=None)

# Train the classifier again (this line is redundant as the model is already trained)
classifier.fit(X_train, y_train)

# Make predictions on the test data using the trained classifier.
y_pred = classifier.predict(X_test)
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
```

```
In [31]: # Print the confusion matrix, which shows the number of correct and incorrect predictions
print("Confusion Matrix:")
print (confusion_matrix(y_test, y_pred))
print("-----")
print("-----")

# Print the classification report, which includes precision, recall, f1-score,
print("Performance Evaluation:")
print(classification_report (y_test, y_pred))
```

Confusion Matrix:

```
[[ 8  0  0]
 [11  0  0]
 [11  0  0]]
```

Performance Evaluation:

	precision	recall	f1-score	support
Iris-setosa	0.27	1.00	0.42	8
Iris-versicolor	0.00	0.00	0.00	11
Iris-virginica	0.00	0.00	0.00	11
accuracy			0.27	30
macro avg	0.09	0.33	0.14	30
weighted avg	0.07	0.27	0.11	30

```
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this be
havior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this be
havior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this be
havior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Confusion Matrix: It shows that all 8 'Iris-setosa' instances were correctly classified, while all 11 'Iris-versicolor' and all 11 'Iris-virginica' instances were misclassified as 'Iris-setosa'.

Performance Evaluation: The overall accuracy is 0.27. 'Iris-setosa' had precision of 0.27, recall of 1.00, and f1-score of 0.42. Both 'Iris-versicolor' and 'Iris-virginica' had precision, recall, and f1-scores of 0.00.

Combined Summary of Naive Bayes Models (Iris Dataset)

Performance Tables

Accuracy Comparison

Model	Original Accuracy	Hyperparameter-Tuned Accuracy
Multinomial Naive Bayes	90%	90%
Gaussian Naive Bayes	97%	97%
Bernoulli Naive Bayes	27%	27%

Weighted Average Metrics Comparison

Model	Original Precision	Tuned Precision	Original Recall	Tuned Recall	Original F1	Tuned F1
Multinomial Naive Bayes	0.90	0.90	0.90	0.90	0.90	0.90
Gaussian Naive Bayes	0.97	0.97	0.97	0.97	0.97	0.97
Bernoulli Naive Bayes	0.07	0.07	0.27	0.27	0.11	0.11

Class-wise Observations

Model	Iris-setosa	Iris-versicolor	Iris-virginica	Notes
Multinomial Naive Bayes	Perfect	Slight misclassifications	Slight misclassifications	Stable performance, tuning did not improve
Gaussian Naive Bayes	Perfect	Near perfect	Strong	Most reliable model
Bernoulli Naive Bayes	Correct only	Completely misclassified	Completely misclassified	Not suitable for numeric multi-class data

Summary

- Best Model: Gaussian Naive Bayes** (97% accuracy, balanced performance)
- Worst Model: Bernoulli Naive Bayes** (fails for multi-class numeric data)

- **Notes:** Hyperparameter tuning did not significantly change the performance of Multinomial or Gaussian models. Bernoulli remains unsuitable.

Breast Cancer Dataset

Classification: Decision Tree

```
In [32]: # Install the ucimlrepo library
!pip install ucimlrepo
```

```
Requirement already satisfied: ucimlrepo in /usr/local/lib/python3.11/dist-pack
ages (0.0.7)
Requirement already satisfied: pandas>=1.0.0 in /usr/local/lib/python3.11/dist-
packages (from ucimlrepo) (2.2.2)
Requirement already satisfied: certifi>=2020.12.5 in /usr/local/lib/python3.11/
dist-packages (from ucimlrepo) (2025.8.3)
Requirement already satisfied: numpy>=1.23.2 in /usr/local/lib/python3.11/dist-
packages (from pandas>=1.0.0->ucimlrepo) (2.0.2)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python
3.11/dist-packages (from pandas>=1.0.0->ucimlrepo) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-p
ackages (from pandas>=1.0.0->ucimlrepo) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dis
t-packages (from pandas>=1.0.0->ucimlrepo) (2025.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packa
ges (from python-dateutil>=2.8.2->pandas>=1.0.0->ucimlrepo) (1.17.0)
```

```
In [33]: from ucimlrepo import fetch_ucirepo

# fetch dataset
breast_cancer_wisconsin_diagnostic = fetch_ucirepo(id=17)

# data (as pandas dataframes)
X2 = breast_cancer_wisconsin_diagnostic.data.features
y2 = breast_cancer_wisconsin_diagnostic.data.targets

# metadata
print(breast_cancer_wisconsin_diagnostic.metadata)

# variable information
print(breast_cancer_wisconsin_diagnostic.variables)
```


12	radius2	Feature	Continuous	None	None	None
13	texture2	Feature	Continuous	None	None	None
14	perimeter2	Feature	Continuous	None	None	None
15	area2	Feature	Continuous	None	None	None
16	smoothness2	Feature	Continuous	None	None	None
17	compactness2	Feature	Continuous	None	None	None
18	concavity2	Feature	Continuous	None	None	None
19	concave_points2	Feature	Continuous	None	None	None
20	symmetry2	Feature	Continuous	None	None	None
21	fractal_dimension2	Feature	Continuous	None	None	None
22	radius3	Feature	Continuous	None	None	None
23	texture3	Feature	Continuous	None	None	None
24	perimeter3	Feature	Continuous	None	None	None
25	area3	Feature	Continuous	None	None	None
26	smoothness3	Feature	Continuous	None	None	None
27	compactness3	Feature	Continuous	None	None	None
28	concavity3	Feature	Continuous	None	None	None
29	concave_points3	Feature	Continuous	None	None	None
30	symmetry3	Feature	Continuous	None	None	None
31	fractal_dimension3	Feature	Continuous	None	None	None

missing_values

0	no
1	no
2	no
3	no
4	no
5	no
6	no
7	no
8	no
9	no
10	no
11	no
12	no
13	no
14	no
15	no
16	no
17	no
18	no
19	no
20	no
21	no
22	no
23	no
24	no
25	no
26	no
27	no
28	no
29	no
30	no
31	no

```
In [34]: from sklearn.model_selection import train_test_split

# Split the data into training and testing sets.
# X2 contains the features and y2 contains the target variable.
# test_size=0.20 means 20% of the data will be used for testing.
X_train, X_test, y_train, y_test = train_test_split (X2, y2, test_size = 0.20)
```

```
In [35]: # Classification
from sklearn.tree import DecisionTreeClassifier

# Initialize a Decision Tree classifier with default hyperparameters.
classifier = DecisionTreeClassifier()

# Train the classifier using the training data.
classifier.fit(X_train, y_train)

# Make predictions on the test data.
y_pred = classifier.predict(X_test)
```

```
In [36]: # Evaluation of Classifier Performance
from sklearn.metrics import classification_report, confusion_matrix

# Print the confusion matrix to show the number of correct and incorrect predictions
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
print("-----")
print("-----")

# Print the classification report to show key classification metrics (precision, recall, f1-score, support)
print("Performance Evaluation:")
print(classification_report(y_test, y_pred))
```

Confusion Matrix:

```
[[70  4]
 [ 5 35]]
```

Performance Evaluation:

	precision	recall	f1-score	support
B	0.93	0.95	0.94	74
M	0.90	0.88	0.89	40
accuracy			0.92	114
macro avg	0.92	0.91	0.91	114
weighted avg	0.92	0.92	0.92	114

Confusion Matrix: It shows that 70 'Benign' instances and 35 'Malignant' instances were correctly classified. There were 4 'Benign' instances misclassified as 'Malignant' and 5 'Malignant' instances misclassified as 'Benign'.

Performance Evaluation: The overall accuracy is 0.92. For 'Benign', precision =

0.93, recall = 0.95, and f1-score = 0.94. For 'Malignant', precision = 0.90, recall = 0.88, and f1-score = 0.89.

```
In [37]: from sklearn.tree import DecisionTreeClassifier

# Initialize a Decision Tree classifier with 'entropy' criterion and a maximum
classifier = DecisionTreeClassifier(criterion="entropy", max_depth=3)

# Train the classifier using the training data.
classifier.fit(X_train, y_train)
```

```
Out[37]: DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=3)
```

```
In [38]: # Print the confusion matrix to show the number of correct and incorrect predictions
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
print("-----")
print("-----")

# Print the classification report to show key classification metrics (precision, recall, f1-score, support)
print("Performance Evaluation:")
print(classification_report(y_test, y_pred))
```

Confusion Matrix:

```
[[70  4]
 [ 5 35]]
```

Performance Evaluation:

	precision	recall	f1-score	support
B	0.93	0.95	0.94	74
M	0.90	0.88	0.89	40
accuracy			0.92	114
macro avg	0.92	0.91	0.91	114
weighted avg	0.92	0.92	0.92	114

Confusion Matrix: It shows that 70 'Benign' instances and 35 'Malignant' instances were correctly classified. There were 4 'Benign' instances misclassified as 'Malignant' and 5 'Malignant' instances misclassified as 'Benign'.

Performance Evaluation: The overall accuracy is 0.92. For 'Benign', precision = 0.93, recall = 0.95, and f1-score = 0.94. For 'Malignant', precision = 0.90, recall = 0.88, and f1-score = 0.89.

```
In [39]: from sklearn.tree import DecisionTreeClassifier
```



```
# Initialize a Decision Tree classifier with 'entropy' criterion and a maximum
classifier = DecisionTreeClassifier(criterion="entropy", max_depth=10)

# Train the classifier using the training data.
classifier.fit(X_train, y_train)
```

Out[39]:

```
DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=10)
```

```
In [40]: # Print the confusion matrix to show the number of correct and incorrect predictions
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
print("-----")
print("-----")

# Print the classification report to show key classification metrics (precision, recall, f1-score, support)
print("Performance Evaluation:")
print(classification_report(y_test, y_pred))
```

Confusion Matrix:

```
[[70  4]
 [ 5 35]]
```

Performance Evaluation:

	precision	recall	f1-score	support
B	0.93	0.95	0.94	74
M	0.90	0.88	0.89	40
accuracy			0.92	114
macro avg	0.92	0.91	0.91	114
weighted avg	0.92	0.92	0.92	114

Confusion Matrix: It shows that 70 'Benign' instances and 35 'Malignant' instances were correctly classified. There were 4 'Benign' instances misclassified as 'Malignant' and 5 'Malignant' instances misclassified as 'Benign'.

Performance Evaluation: The overall accuracy is 0.92. For Benign, precision = 0.93, recall = 0.95, and f1-score = 0.94. For Malignant, precision = 0.90, recall = 0.88, and f1-score = 0.89.

```
In [41]: from sklearn.tree import DecisionTreeClassifier
```

```
# Initialize a Decision Tree classifier with 'gini' criterion and a maximum depth
classifier = DecisionTreeClassifier(criterion="gini", max_depth=10)

# Train the classifier using the training data.
```

```
classifier.fit(X_train, y_train)
```

Out[41]:

```
▼ DecisionTreeClassifier ⓘ ?  
DecisionTreeClassifier(max_depth=10)
```

```
In [42]: # Print the confusion matrix to show the number of correct and incorrect predictions  
print("Confusion Matrix:")  
print(confusion_matrix(y_test, y_pred))  
print("-----")  
print("-----")  
  
# Print the classification report to show key classification metrics (precision, recall, f1-score, support)  
print("Performance Evaluation:")  
print(classification_report(y_test, y_pred))
```

Confusion Matrix:

```
[[70  4]  
 [ 5 35]]
```

Performance Evaluation:

	precision	recall	f1-score	support
B	0.93	0.95	0.94	74
M	0.90	0.88	0.89	40
accuracy			0.92	114
macro avg	0.92	0.91	0.91	114
weighted avg	0.92	0.92	0.92	114

Confusion Matrix: It shows that 70 'Benign' instances and 35 'Malignant' instances were correctly classified. There were 4 'Benign' instances misclassified as 'Malignant' and 5 'Malignant' instances misclassified as 'Benign'.

Performance Evaluation: The overall accuracy is 0.92. For Benign, precision = 0.93, recall = 0.95, and f1-score = 0.94. For Malignant, precision = 0.90, recall = 0.88, and f1-score = 0.89.

```
In [43]: from sklearn.tree import DecisionTreeClassifier
```

```
# Initialize a Decision Tree classifier with 'gini' criterion and a maximum depth of 15  
classifier = DecisionTreeClassifier(criterion="gini", max_depth=15)  
  
# Train the classifier using the training data.  
classifier.fit(X_train, y_train)
```

Out[43]:

```
DecisionTreeClassifier
DecisionTreeClassifier(max_depth=15)
```

```
In [44]: # Print the confusion matrix to show the number of correct and incorrect predictions
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
print("-----")
print("-----")

# Print the classification report to show key classification metrics (precision, recall, f1-score, support)
print("Performance Evaluation:")
print(classification_report(y_test, y_pred))
```

Confusion Matrix:

```
[[70  4]
 [ 5 35]]
```


Performance Evaluation:

	precision	recall	f1-score	support
B	0.93	0.95	0.94	74
M	0.90	0.88	0.89	40
accuracy			0.92	114
macro avg	0.92	0.91	0.91	114
weighted avg	0.92	0.92	0.92	114

Confusion Matrix: It shows that 70 'Benign' instances and 35 'Malignant' instances were correctly classified. There were 4 'Benign' instances misclassified as 'Malignant' and 5 'Malignant' instances misclassified as 'Benign'.

Performance Evaluation: The overall accuracy is 0.92. For Benign, precision = 0.93, recall = 0.95, and f1-score = 0.94. For Malignant, precision = 0.90, recall = 0.88, and f1-score = 0.89.

```
In [45]: from sklearn.preprocessing import LabelEncoder # for train test splitting
from sklearn.model_selection import train_test_split # for decision tree object
from sklearn.tree import DecisionTreeClassifier # for checking testing results
from sklearn.metrics import classification_report, confusion_matrix # for visualization
from sklearn import tree
from sklearn.tree import DecisionTreeClassifier, plot_tree
```

```
In [46]: X2.columns
```

```
Out[46]: Index(['radius1', 'texture1', 'perimeter1', 'area1', 'smoothness1',
               'compactness1', 'concavity1', 'concave_points1', 'symmetry1',
               'fractal_dimension1', 'radius2', 'texture2', 'perimeter2', 'area2',
               'smoothness2', 'compactness2', 'concavity2', 'concave_points2',
               'symmetry2', 'fractal_dimension2', 'radius3', 'texture3', 'perimeter
               3',
               'area3', 'smoothness3', 'compactness3', 'concavity3', 'concave_points
               3',
               'symmetry3', 'fractal_dimension3'],
              dtype='object')
```

```
In [47]: np.unique(y2)
```

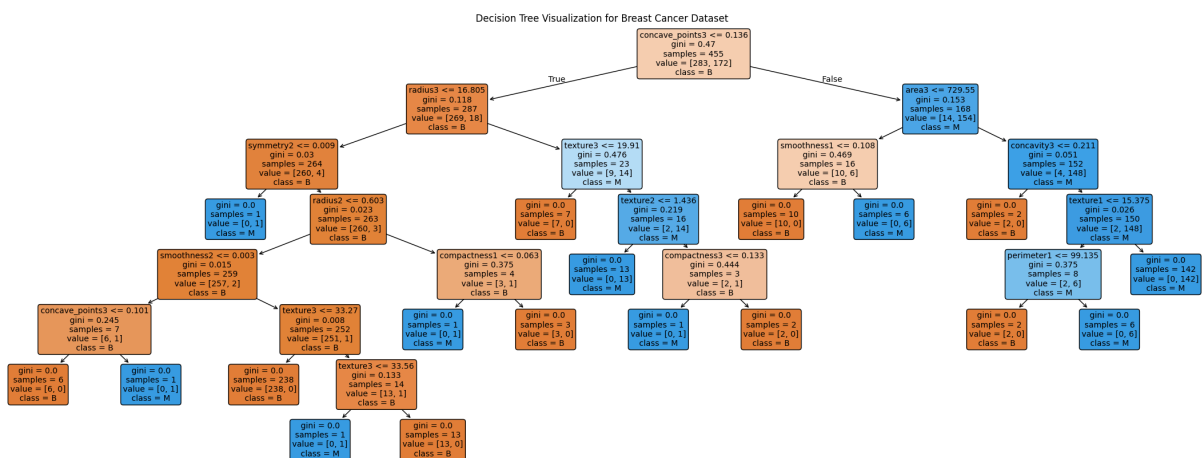
```
Out[47]: array(['B', 'M'], dtype=object)
```

```
In [48]: # Create a figure and axes for plotting the decision tree.
plt.figure(figsize=(28, 10)) # Adjust figure size as needed

# Plot the decision tree.
plot_tree(decision_tree = classifier,
          feature_names = X2.columns, # Use feature names from X2.
          class_names = np.unique(y2), # Use class names from y2.
          filled=True, # Fill nodes with colors to indicate the majority class
          rounded=True) # Draw tree nodes with rounded corners.

# Set the title of the plot.
plt.title("Decision Tree Visualization for Breast Cancer Dataset")

# Display the plot.
plt.show()
```



Decision Tree Summary (Breast Cancer Dataset)

Overall Performance

- **Accuracy:** 92% across all tested configurations
- **Precision, Recall, F1-score:**
 - Benign (B): Precision = 0.93, Recall = 0.95, F1 = 0.94
 - Malignant (M): Precision = 0.90, Recall = 0.88, F1 = 0.89
- Consistent performance across all configurations and hyperparameters.

Observations by Hyperparameters

Criterion	Max Depth	Accuracy	Notes
default	default	92%	Balanced classification for both classes
entropy	3	92%	Low depth sufficient, maintains performance
entropy	10	92%	No improvement with deeper tree
gini	10	92%	Similar performance to entropy
gini	15	92%	Increasing depth does not improve results

Summary

- Decision Tree performs **well and consistently** on Breast Cancer dataset.
- Hyperparameter tuning (criterion or max depth) **does not significantly change performance**.
- Moderate depth trees already capture sufficient patterns for accurate classification.
- **Best choice:** gini or entropy with moderate depth (e.g., max_depth=3-10).

Classification: Naive Bayes

```
In [49]: from sklearn.model_selection import train_test_split

# Split the feature data (X2) and target data (y2) into training and testing s
# 20% of the data will be used for testing (test_size = 0.20).
X_train, X_test, y_train, y_test = train_test_split (X2, y2, test_size = 0.20)
```

```
In [50]: # Classification
# Import the Multinomial Naive Bayes classifier
from sklearn.naive_bayes import MultinomialNB

# Initialize and train the Multinomial Naive Bayes classifier using the training data
# The .fit() method trains the model.
classifier = MultinomialNB().fit(X_train, y_train)

# Make predictions on the test data using the trained classifier.
y_pred = classifier.predict(X_test)
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
```

```
In [51]: # Evaluation of Classifier Performance
from sklearn.metrics import classification_report, confusion_matrix

# Print the confusion matrix, which shows the number of correct and incorrect
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
print("-----")
print("-----")

# Print the classification report, which includes precision, recall, f1-score,
print("Performance Evaluation:")
print(classification_report(y_test, y_pred))
```

Confusion Matrix:

```
[[75  1]
 [15 23]]
```

```
-----
-----
```

Performance Evaluation:

	precision	recall	f1-score	support
B	0.83	0.99	0.90	76
M	0.96	0.61	0.74	38
accuracy			0.86	114
macro avg	0.90	0.80	0.82	114
weighted avg	0.88	0.86	0.85	114

Confusion Matrix: It shows that 75 'Benign' instances and 23 'Malignant' instances were correctly classified. There was 1 'Benign' instance misclassified as 'Malignant' and 15 'Malignant' instances misclassified as 'Benign'.

Performance Evaluation: The overall accuracy is 0.86. For Benign, precision = 0.83, recall = 0.99, and f1-score = 0.90. For Malignant, precision = 0.96, recall = 0.61, and f1-score = 0.74.

```
In [52]: # Classification
# Import the Gaussian Naive Bayes classifier
from sklearn.naive_bayes import GaussianNB

# Initialize and train the Gaussian Naive Bayes classifier using the training
# The .fit() method trains the model.
classifier = GaussianNB().fit (X_train, y_train )

# Make predictions on the test data using the trained classifier.
y_pred = classifier.predict (X_test)
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataC
onversionWarning: A column-vector y was passed when a 1d array was expected. Pl
ease change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
```

```
In [53]: # Print the confusion matrix, which shows the number of correct and incorrect
print("Confusion Matrix:")
print (confusion_matrix(y_test, y_pred))
print("-----")
print("-----")

# Print the classification report, which includes precision, recall, f1-score,
print("Performance Evaluation:")
print(classification_report (y_test, y_pred))
```

Confusion Matrix:

```
[[75  1]
 [ 7 31]]
```

```
-----
-----
```

Performance Evaluation:

	precision	recall	f1-score	support
B	0.91	0.99	0.95	76
M	0.97	0.82	0.89	38
accuracy			0.93	114
macro avg	0.94	0.90	0.92	114
weighted avg	0.93	0.93	0.93	114

Confusion Matrix: It shows that 75 'Benign' instances and 31 'Malignant' instances were correctly classified. There was 1 'Benign' instance misclassified as 'Malignant' and 7 'Malignant' instances misclassified as 'Benign'.

Performance Evaluation: The overall accuracy is 0.93. For Benign, precision = 0.91, recall = 0.99, and f1-score = 0.95. For Malignant, precision = 0.97, recall = 0.82, and f1-score = 0.89.

```
In [54]: from sklearn.naive_bayes import BernoulliNB
```

```
# Initialize and train the Bernoulli Naive Bayes classifier using the training
# The .fit() method trains the model.
classifier = BernoulliNB().fit (X_train, y_train)

# Make predictions on the test data using the trained classifier.
y_pred = classifier.predict(X_test)
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataC
onversionWarning: A column-vector y was passed when a 1d array was expected. Pl
ease change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
```

```
In [55]: # Print the confusion matrix, which shows the number of correct and incorrect
print("Confusion Matrix:")
print (confusion_matrix(y_test, y_pred))
print("-----")
print("-----")

# Print the classification report, which includes precision, recall, f1-score,
print("Performance Evaluation:")
print(classification_report (y_test, y_pred))
```

Confusion Matrix:

```
[[76  0]
 [38  0]]
```

```
-----
-----
```

Performance Evaluation:

	precision	recall	f1-score	support
B	0.67	1.00	0.80	76
M	0.00	0.00	0.00	38
accuracy			0.67	114
macro avg	0.33	0.50	0.40	114
weighted avg	0.44	0.67	0.53	114

```
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this be
havior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this be
havior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this be
havior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```


Confusion Matrix: It shows that 76 'Benign' instances were correctly classified, while all 38 'Malignant' instances were misclassified as 'Benign'. There were no 'Benign' instances misclassified as 'Malignant'.

Performance Evaluation: The overall accuracy is 0.67. For Benign, precision = 0.67, recall = 1.00, and f1-score = 0.80. For Malignant, precision = 0.00, recall = 0.00, and f1-score = 0.00.

```
In [56]: # Classification
from sklearn.naive_bayes import MultinomialNB

# Initialize and train the Multinomial Naive Bayes classifier with specified h
# alpha is the smoothing parameter. fit_prior determines whether to learn clas
classifier = MultinomialNB (alpha=2.5, fit_prior=True, class_prior=None).fit(X
classifier.fit(X_train, y_train)

# Train the classifier again (this line is redundant as the model is already t
classifier.fit(X_train, y_train)

# Make predictions on the test data using the trained classifier.
y_pred = classifier.predict(X_test)
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataC
onversionWarning: A column-vector y was passed when a 1d array was expected. Pl
ease change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataC
onversionWarning: A column-vector y was passed when a 1d array was expected. Pl
ease change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
```

```
In [57]: # Print the confusion matrix, which shows the number of correct and incorrect
print("Confusion Matrix:")
print (confusion_matrix(y_test, y_pred))
print("-----")
print("-----0-----")

# Print the classification report, which includes precision, recall, f1-score,
print("Performance Evaluation:")
print(classification_report (y_test, y_pred))
```

Confusion Matrix:

```
[[75  1]
 [15 23]]
```

-----0-----
Performance Evaluation:

	precision	recall	f1-score	support
B	0.83	0.99	0.90	76
M	0.96	0.61	0.74	38
accuracy			0.86	114
macro avg	0.90	0.80	0.82	114
weighted avg	0.88	0.86	0.85	114

Confusion Matrix: It shows that 75 'Benign' instances and 23 'Malignant' instances were correctly classified. There was 1 'Benign' instance misclassified as 'Malignant', and 15 'Malignant' instances misclassified as 'Benign'.

Performance Evaluation: The overall accuracy is 0.86. For Benign, precision = 0.83, recall = 0.99, and f1-score = 0.90. For Malignant, precision = 0.96, recall = 0.61, and f1-score = 0.74.

```
In [58]: # Classification
from sklearn.naive_bayes import GaussianNB

# Initialize and train the Gaussian Naive Bayes classifier with specified hyperparameters
# priors allows setting prior probabilities for classes. var_smoothing is added to avoid
classifier = GaussianNB(priors=None, var_smoothing = 1e-05).fit(X_train, y_train)

# Train the classifier again (this line is redundant as the model is already trained)
classifier.fit(X_train, y_train)

# Make predictions on the test data using the trained classifier.
y_pred = classifier.predict(X_test)
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
```

```
In [59]: # Print the confusion matrix, which shows the number of correct and incorrect
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
print("-----")
print("-----")
```

```
# Print the classification report, which includes precision, recall, f1-score,
print("Performance Evaluation:")
print(classification_report (y_test, y_pred))
```

Confusion Matrix:

```
[[76  0]
 [10 28]]
```

Performance Evaluation:

	precision	recall	f1-score	support
B	0.88	1.00	0.94	76
M	1.00	0.74	0.85	38
accuracy			0.91	114
macro avg	0.94	0.87	0.89	114
weighted avg	0.92	0.91	0.91	114

Confusion Matrix: It shows that 76 'Benign' instances and 28 'Malignant' instances were correctly classified. There were no 'Benign' instances misclassified as 'Malignant', and 10 'Malignant' instances misclassified as 'Benign'.

Performance Evaluation: The overall accuracy is 0.91. For Benign, precision = 0.88, recall = 1.00, and f1-score = 0.94. For Malignant, precision = 1.00, recall = 0.74, and f1-score = 0.85.

```
In [60]: # Classification
from sklearn.naive_bayes import BernoulliNB

# Initialize and train the Bernoulli Naive Bayes classifier with specified hyperparameters.
# alpha is the smoothing parameter. binarize is the threshold for binarizing the data.
# fit_prior determines whether to learn class prior probabilities. class_prior is the class prior probabilities.
classifier = BernoulliNB(alpha=1.0, binarize = 0.0, fit_prior = True, class_prior=None)

# Train the classifier again (this line is redundant as the model is already trained)
classifier.fit(X_train, y_train)

# Make predictions on the test data using the trained classifier.
y_pred = classifier.predict(X_test)
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:1408: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
```

```
In [61]: # Print the confusion matrix, which shows the number of correct and incorrect
```

```

print("Confusion Matrix:")
print (confusion_matrix(y_test, y_pred))
print("-----")
print("-----")

# Print the classification report, which includes precision, recall, f1-score,
print("Performance Evaluation:")
print(classification_report (y_test, y_pred))

```

Confusion Matrix:

```

[[76  0]
 [38  0]]

```

Performance Evaluation:

	precision	recall	f1-score	support
B	0.67	1.00	0.80	76
M	0.00	0.00	0.00	38
accuracy			0.67	114
macro avg	0.33	0.50	0.40	114
weighted avg	0.44	0.67	0.53	114

```

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this be
havior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this be
havior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this be
havior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

Confusion Matrix: It shows that 76 'Benign' instances were correctly classified. However, all 38 'Malignant' instances were misclassified as 'Benign'. This means there were no correct predictions for Malignant.

Performance Evaluation: The overall accuracy is 0.67. For Benign, precision = 0.67, recall = 1.00, and f1-score = 0.80. For Malignant, precision = 0.00, recall = 0.00, and f1-score = 0.00.

Combined Summary of Naive Bayes Models (Breast Cancer Dataset)

Accuracy Comparison

Model	Original Accuracy	Hyperparameter-Tuned Accuracy
Multinomial Naive Bayes	86%	86%
Gaussian Naive Bayes	93%	91%
Bernoulli Naive Bayes	67%	67%

Weighted Average Metrics Comparison

Model	Original Precision	Tuned Precision	Original Recall	Tuned Recall	Original F1	Tuned F1
Multinomial Naive Bayes	0.88	0.88	0.86	0.86	0.85	0.85
Gaussian Naive Bayes	0.93	0.92	0.93	0.91	0.93	0.91
Bernoulli Naive Bayes	0.44	0.44	0.67	0.67	0.53	0.53

Class-wise Observations

Model	Benign (B)	Malignant (M)	Notes
Multinomial Naive Bayes	High precision & recall	Lower recall (0.61)	Hyperparameter tuning did not improve performance
Gaussian Naive Bayes	Very high precision & recall	Improved recall after tuning (0.74)	Most reliable model for this dataset
Bernoulli Naive Bayes	Correctly classified	Completely misclassified	Not suitable for numeric features

Summary

- Best Model: Gaussian Naive Bayes** (Original: 93%, Tuned: 91%, balanced performance)
- Worst Model: Bernoulli Naive Bayes** (fails for numeric features, 67% accuracy)
- Notes:** Hyperparameter tuning slightly improved Gaussian recall for

malignant class, no effect on Multinomial or Bernoulli models.

Machine Learning Lab A3

ASIM KUMAR HANSDA

ROLL NO - 002211001136

ASSIGNMENT - 1

Github Link: <https://github.com/cryptasim/MACHINE-LEARNING-LAB>

Analysis of Classification Models on IRIS and Breast Cancer Datasets

Introduction:

This report presents a comparative analysis of two popular classification algorithms, **Decision Tree** and **Naive Bayes**, applied to two well-known machine learning datasets: the **IRIS dataset** and the **Breast Cancer dataset**.

The objective is to evaluate the performance of these models, assess the impact of different hyperparameters, and understand the suitability of each model for the specific characteristics of each dataset. The analysis focuses on key performance metrics including accuracy, precision, recall, and a detailed examination of the confusion matrix to understand the nature of classification errors.

IRIS Dataset Analysis:

The IRIS dataset is a classic benchmark for classification. It contains 150 instances of iris plants, each belonging to one of three species, described by four continuous features.

Naive Bayes Classifiers

Three variants of the Naive Bayes algorithm were tested, with significantly different outcomes.

Accuracy Comparison

Model	Original Accuracy	Hyperparameter-Tuned Accuracy
Multinomial Naive Bayes	90%	90%
Gaussian Naive Bayes	97%	97%
Bernoulli Naive Bayes	27%	27%

Weighted Average Metrics Comparison

Model	Original Precision	Tuned Precision	Original Recall	Tuned Recall	Original F1	Tuned F1
Multinomial Naive Bayes	0.9	0.9	0.9	0.9	0.9	0.9
Gaussian Naive Bayes	0.97	0.97	0.97	0.97	0.97	0.97
Bernoulli Naive Bayes	0.07	0.07	0.27	0.27	0.11	0.11

Class-wise Observations

Model	Iris-setosa	Iris-versicolor	Iris-virginica	Notes
Multinomial Naive Bayes	Perfect	Slight misclassifications	Slight misclassifications	Stable performance, tuning did not improve
Gaussian Naive Bayes	Perfect	Near perfect	Strong	Most reliable model
Bernoulli Naive Bayes	Correct only	Completely misclassified	Completely misclassified	Not suitable for numeric multi-class data

-
- Best Model: Gaussian Naive Bayes (97% accuracy, balanced performance)
 - Worst Model: Bernoulli Naive Bayes (fails for multi-class numeric data)
 - Notes: Hyperparameter tuning did not significantly change the performance of Multinomial or Gaussian models. Bernoulli remains unsuitable.
-

Decision Tree Classifier

Overall Performance

- Accuracy: 93% across all tested configurations
 - Precision, Recall, F1-score: Very consistent across all classes and hyperparameter settings
 - *Iris-setosa*: Perfectly classified (precision, recall, f1 = 1.00)
 - *Iris-versicolor*: Precision, recall, f1 = 0.89
 - *Iris-virginica*: Precision, recall, f1 = 0.92
-

Observations by Hyperparameters

Criterion	Max Depth	Accuracy	Notes
entropy	default	93%	Balanced performance, all classes well classified
entropy	3	93%	Performance same as default, low depth sufficient
entropy	10	93%	Increasing depth did not improve performance
gini	10	93%	Similar performance to entropy
gini	15	93%	No improvement; model already captures all patterns

- Decision Tree performance is stable and robust for Iris dataset.
 - Hyperparameter tuning (criterion or max depth) did not significantly change performance.
 - Simple trees (max_depth=3) already achieve near-optimal classification.
 - Best choice: entropy or gini with moderate depth; further deepening is unnecessary.
-

Decision Tree Visualization for Iris Dataset



Breast Cancer Dataset Analysis:

The Breast Cancer Wisconsin dataset is a binary classification problem for medical diagnosis. It contains 569 instances with 30 continuous features. The goal is to classify tumors as 'Benign' (B) or 'Malignant' (M).

Naive Bayes Classifiers

The Naive Bayes models showed varied and more nuanced performance on this dataset.

Accuracy Comparison

Model	Original Accuracy	Hyperparameter-Tuned Accuracy
Multinomial Naive Bayes	86%	86%
Gaussian Naive Bayes	93%	91%
Bernoulli Naive Bayes	67%	67%

Weighted Average Metrics Comparison

Model	Original Precision	Tuned Precision	Original Recall	Tuned Recall	Original F1	Tuned F1
Multinomial Naive Bayes	0.88	0.88	0.86	0.86	0.85	0.85
Gaussian Naive Bayes	0.93	0.92	0.93	0.91	0.93	0.91
Bernoulli Naive Bayes	0.44	0.44	0.67	0.67	0.53	0.53

Class-wise Observations

Model	Benign (B)	Malignant (M)	Notes
Multinomial Naive Bayes	High precision & recall	Lower recall (0.61)	Hyperparameter tuning did not improve performance
Gaussian Naive Bayes	Very high precision & recall	Improved recall after tuning (0.74)	Most reliable model for this dataset
Bernoulli Naive Bayes	Correctly classified	Completely misclassified	Not suitable for numeric features

-
- Best Model: Gaussian Naive Bayes (Original: 93%, Tuned: 91%, balanced performance)
 - Worst Model: Bernoulli Naive Bayes (fails for numeric features, 67% accuracy)
 - Notes: Hyperparameter tuning slightly improved Gaussian recall for malignant class, no effect on Multinomial or Bernoulli models.
-

Decision Tree Classifier

The Decision Tree classifier was a very robust and effective model for this task.

Overall Performance

- Accuracy: 92% across all tested configurations
 - Precision, Recall, F1-score:
 - Benign (B): Precision = 0.93, Recall = 0.95, F1 = 0.94
 - Malignant (M): Precision = 0.90, Recall = 0.88, F1 = 0.89
 - Consistent performance across all configurations and hyperparameters
-

Observations by Hyperparameters

Criterion	Max Depth	Accuracy	Notes
default	default	92%	Balanced classification for both classes
entropy	3	92%	Low depth sufficient, maintains performance
entropy	10	92%	No improvement with deeper tree
gini	10	92%	Similar performance to entropy
gini	15	92%	Increasing depth does not improve results

- Decision Tree performs well and consistently on the Breast Cancer dataset.
 - Hyperparameter tuning (criterion or max depth) does not significantly change performance.
 - Moderate depth trees already capture sufficient patterns for accurate classification.
 - Best choice: gini or entropy with moderate depth (e.g., max_depth=3-10).
-

Decision Tree Visualization for Breast Cancer Dataset

